

### Review Group 9 – HW3

Your code looks very nice and includes the Replay Buffer, the DQ network and appropriate training algorithm. The comments and additional documentation make it easier to understand and provide a nice overview. For additional readability, you could maybe label the graph axis, as this is the first thing people will look at. Training always takes a lot of time, but you could have trained on more than 10 episodes, the results don't really become apparent with so few iterations. Good job with the implementation.